

8106hw4

Ze Li

Contents

Problem 1	1
Random Forest	4
Random Forest - Variable Importance	5
Random Forest - Test Error	7
Boosting	7
Boosting - Variable Importance	8
Boosting - Test Error	10
Problem 2	10
Classification Tree	11
lse	13
Boosting	14
Test Performance	16

```
library(ISLR)
library(caret)
library(mgcv)
library(tidymodels)
library(rpart)
library(rpart.plot)
library(randomForest)
library(ranger)
library(gbm)
library(pROC)
```

Problem 1

```
college=read.csv("/Users/zeze/Library/Mobile Documents/com~apple~CloudDocs/2024/24S BIST P8106 DS II/hw
indexTrain <- createDataPartition(y = college$Outstate, p = 0.8, list = FALSE)
train <- college[indexTrain, ][-1]
test <- college[-indexTrain, ][-1]
train <- na.omit(train)
```

```
test <- na.omit(test)
head(train)
```

```
##   Apps Accept Enroll Top10perc Top25perc F.Undergrad P.Undergrad Outstate
## 1 1660  1232   721      23        52      2885      537      7440
## 2 2186  1924   512      16        29      2683     1227     12280
## 3 1428  1097   336      22        50      1036       99     11250
## 5  193   146    55      16        44       249      869     7560
## 6  587   479   158      38        62       678       41     13500
## 7  353   340   103      17        45       416      230     13290
##   Room.Board Books Personal PhD Terminal S.F.Ratio perc.alumni Expend Grad.Rate
## 1      3300   450    2200  70      78      18.1      12    7041      60
## 2      6450   750    1500  29      30      12.2      16   10527      56
## 3      3750   400    1165  53      66      12.9      30    8735      54
## 5      4120   800    1500  76      72      11.9       2   10922      15
## 6      3335   500     675  67      73       9.4      11    9727      55
## 7      5720   500    1500  90      93      11.5      26    8861      63
```

```
# matrix of predictors
x_train <- model.matrix(Outstate ~ ., train)[, -1]
head(x_train)
```

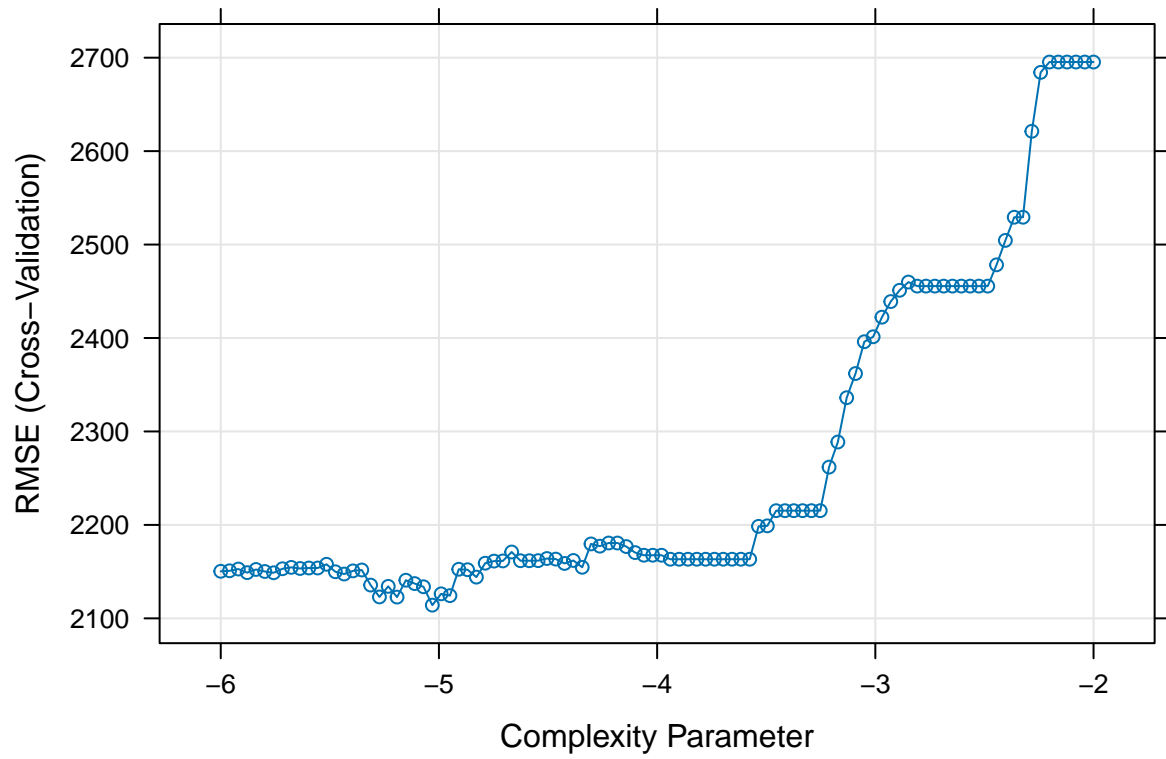
```
##   Apps Accept Enroll Top10perc Top25perc F.Undergrad P.Undergrad Room.Board
## 1 1660  1232   721      23        52      2885      537      3300
## 2 2186  1924   512      16        29      2683     1227     6450
## 3 1428  1097   336      22        50      1036       99     3750
## 5  193   146    55      16        44       249      869     4120
## 6  587   479   158      38        62       678       41     3335
## 7  353   340   103      17        45       416      230     5720
##   Books Personal PhD Terminal S.F.Ratio perc.alumni Expend Grad.Rate
## 1   450    2200  70      78      18.1      12    7041      60
## 2   750    1500  29      30      12.2      16   10527      56
## 3   400    1165  53      66      12.9      30    8735      54
## 5   800    1500  76      72      11.9       2   10922      15
## 6   500     675  67      73       9.4      11    9727      55
## 7   500    1500  90      93      11.5      26    8861      63
```

```
# vector of response
y_train <- train$Outstate
# matrix of predictors
x_test <- model.matrix(Outstate ~ ., test)[, -1]
# vector of response
y_test <- test$Outstate
```

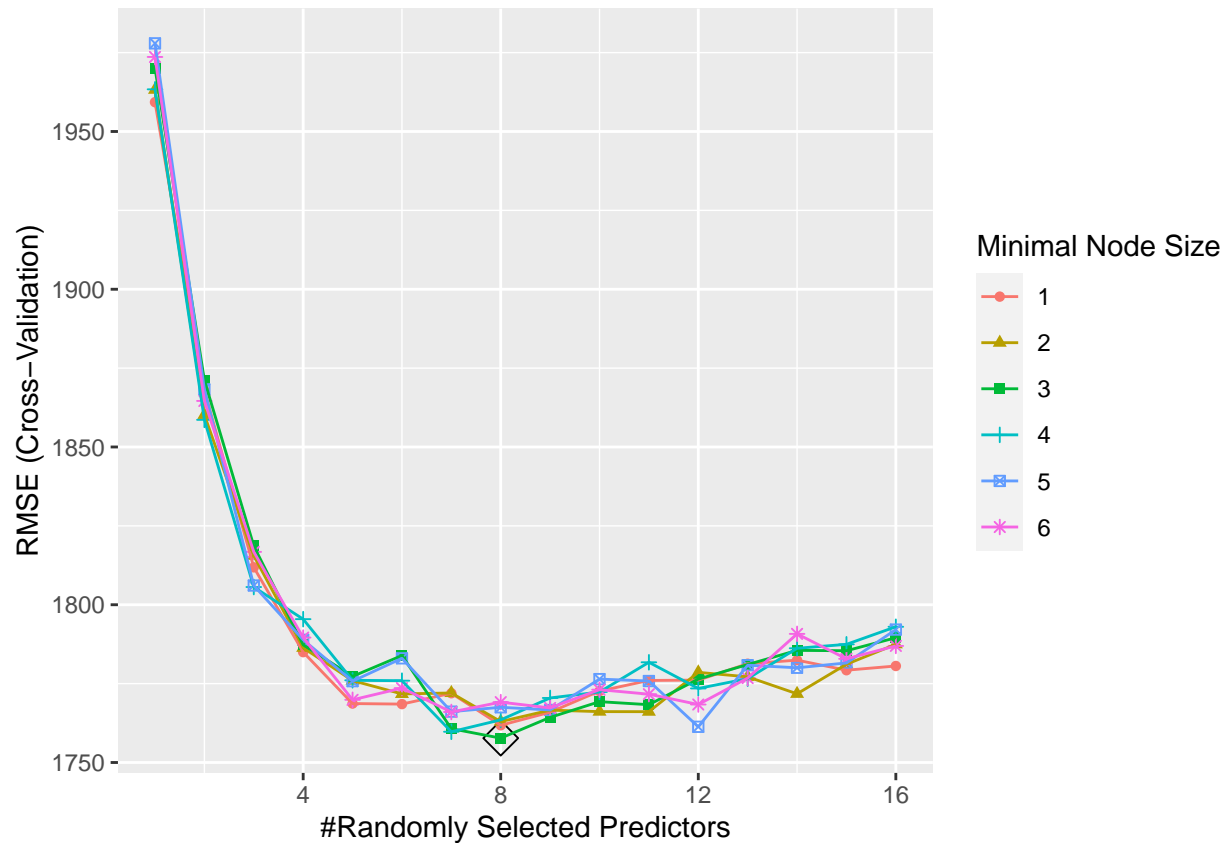
(a) Build a regression tree on the training data to predict the response. Create a plot of the tree.

```
ctrl <- trainControl(method = "cv")
set.seed(1)
rpart.fit <- train(Outstate ~ ., train,
  method = "rpart",
  tuneGrid = data.frame(cp = exp(seq(-6, -2, length = 100))),
```

```
trControl = ctrl)  
plot(rpart.fit, xTrans = log)
```



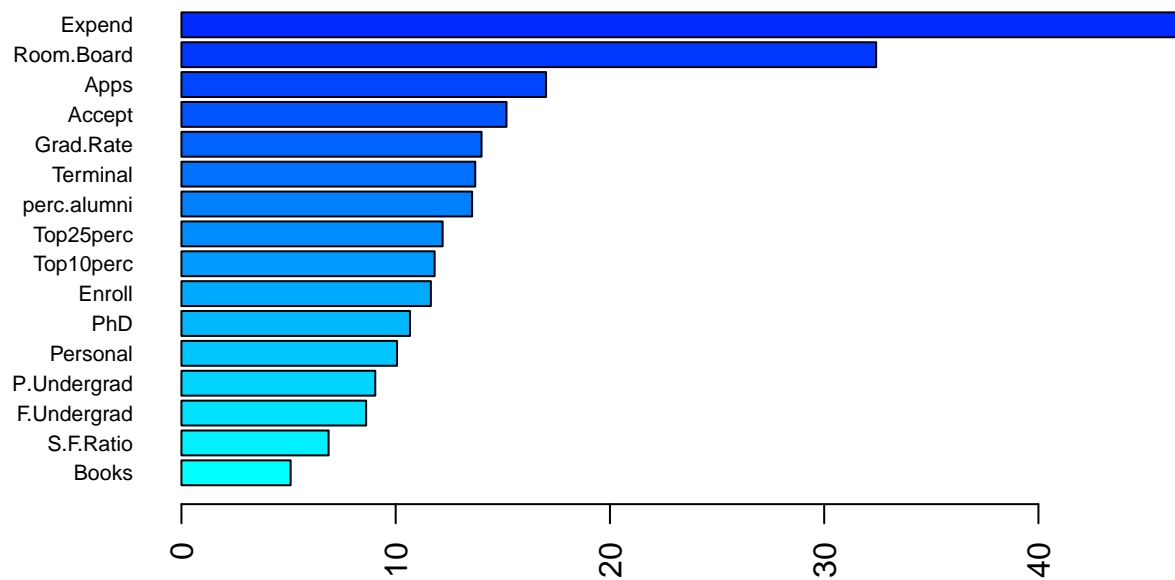
```
rpart.plot(rpart.fit$finalModel)
```

Random Forest - Variable Importance

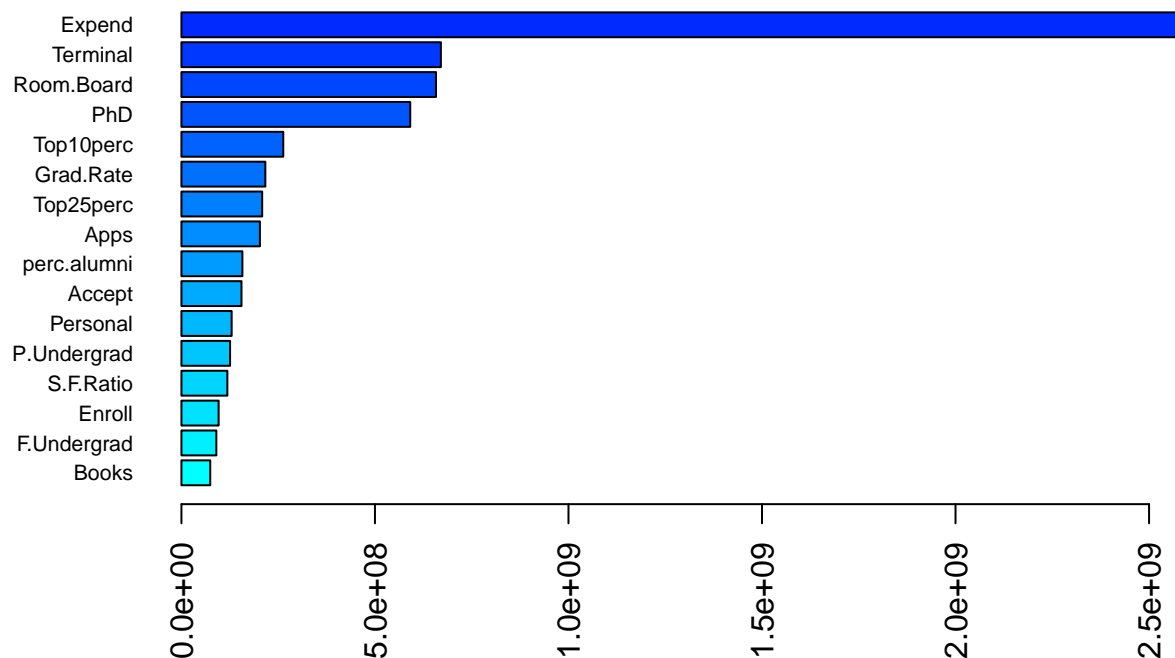
```
set.seed(1)
rf2.final.per <- ranger(Outstate ~ . ,
  data = train,
  mtry = rf.fit$bestTune[[1]],
  splitrule = "variance",
  min.node.size = rf.fit$bestTune[[3]],
  importance = "permutation",
  scale.permutation.importance = TRUE)

barplot(sort(ranger::importance(rf2.final.per), decreasing = FALSE),
  las = 2, horiz = TRUE, cex.names = 0.7,
  col = colorRampPalette(colors = c("cyan", "blue"))(19))
```



```
set.seed(1)
rf2.final.imp <- ranger(Outstate ~ . ,
                        data = train,
                        mtry = rf.fit$bestTune[[1]],
                        splitrule = "variance",
                        min.node.size = rf.fit$bestTune[[3]],
                        importance = "impurity")

barplot(sort(ranger::importance(rf2.final.imp), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("cyan", "blue"))(19))
```



Random Forest - Test Error

```
pred.rf <- predict(rf.fit, newdata = test)
RMSE(pred.rf, y_test)
```

```
## [1] 1608.216
```

The two bar plots illustrate the variable importance derived from a random forest model, with the first plot representing permutation importance and the second reflecting impurity importance. In both metrics, 'Expend' stands out as the most influential predictor, indicating that the amount spent per student is a key factor in predicting the response variable 'Outstate'. Academic-related variables such as 'Room.Board', 'Terminal', 'PhD', and 'Top10perc' also rank highly across both importance measures, underscoring the relevance of financial and educational quality factors in the model's predictions.

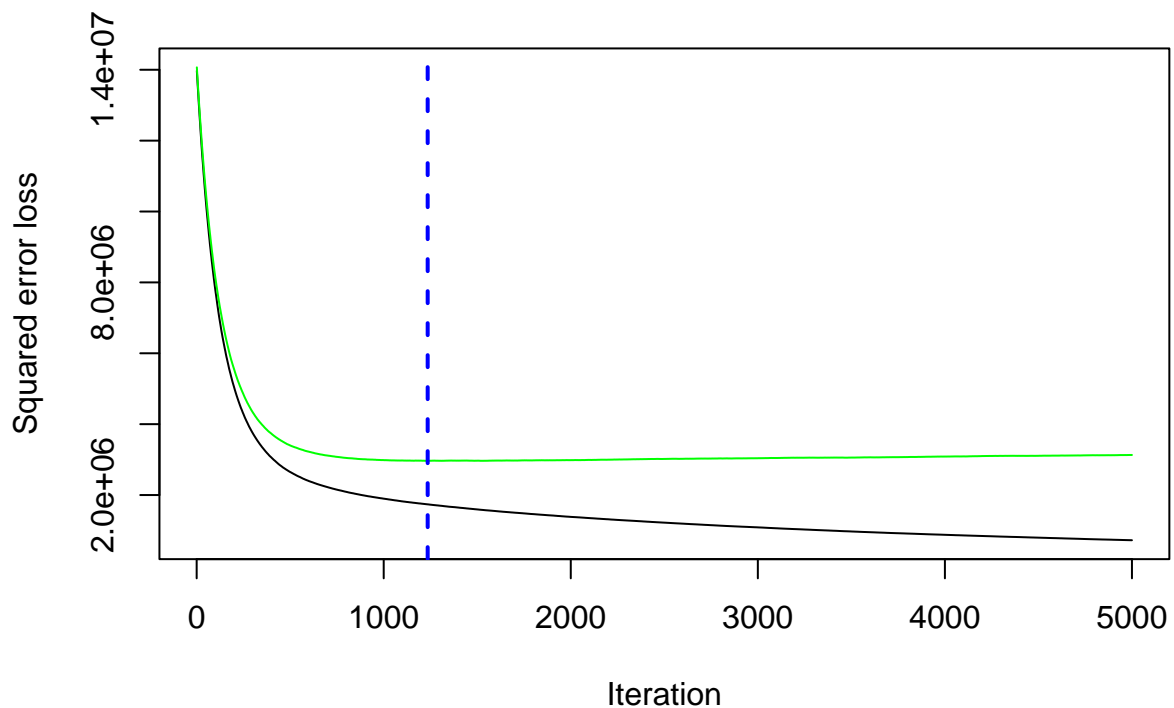
The test error is 1608.215614.

(c) Perform boosting on the training data. Report the variable importance and the test error.

Boosting

```
# We first fit a gradient boosting model with Gaussian loss function
set.seed(1)
```

```
bst <- gbm(Outstate ~ . ,
  data = train,
  distribution = "gaussian",
  n.trees = 5000,
  interaction.depth = 3,
  shrinkage = 0.005,
  cv.folds = 10,
  n.cores = 2)
# We plot loss function as a result of number of trees added to the ensemble
gbm.perf(bst, method = "cv")
```



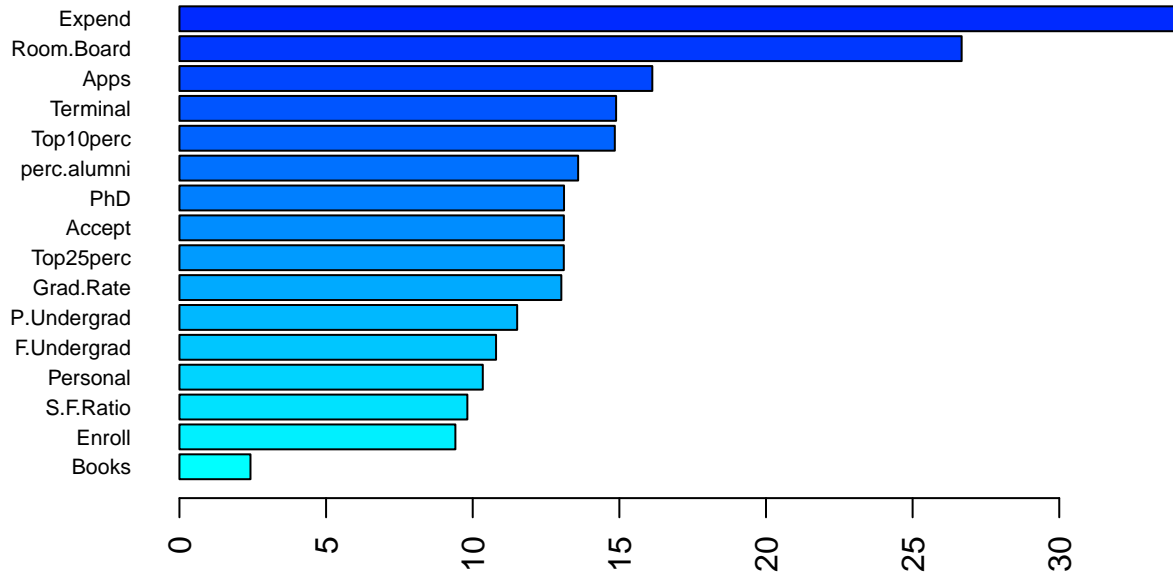
```
## [1] 1235
```

Boosting - Variable Importance

```
set.seed(1)
gbm.final.per <- ranger(Outstate ~ . ,
  data = train,
  mtry = bst$bestTune[[1]],
  splitrule = "variance",
  min.node.size = bst$bestTune[[3]],
  importance = "permutation",
  scale.permutation.importance = TRUE)
```

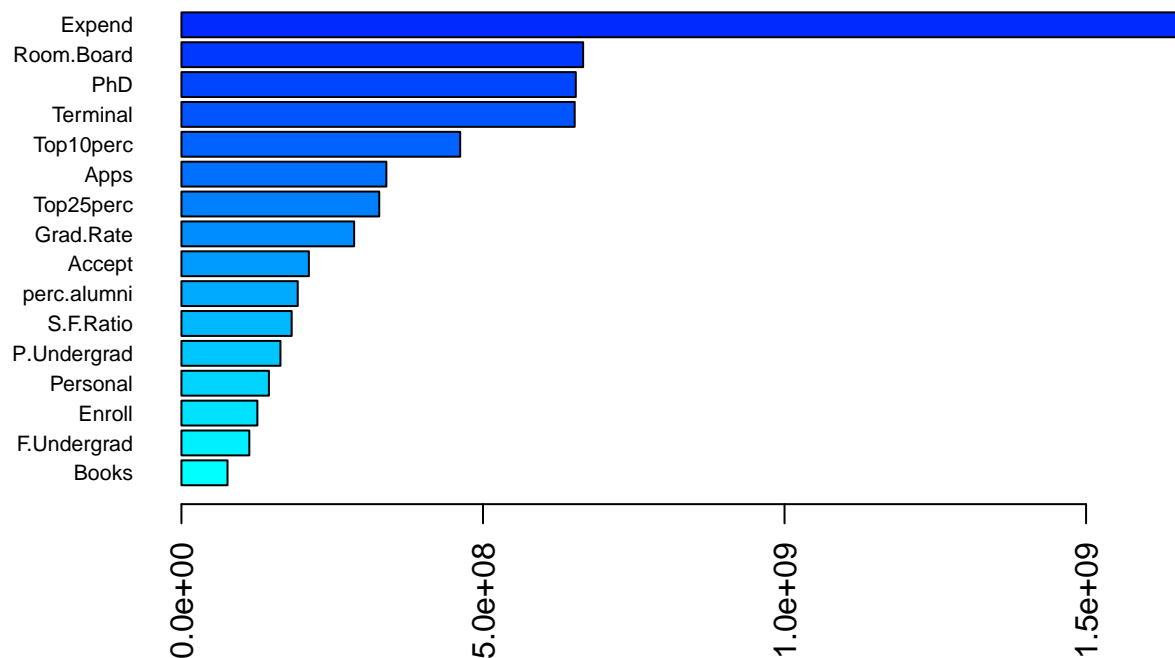


```
barplot(sort(ranger::importance(gbm.final.per), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("cyan", "blue"))(19))
```



```
set.seed(1)
gbm.final.imp <- ranger(Outstate ~ . ,
                        data = train,
                        mtry = bst$bestTune[[1]],
                        splitrule = "variance",
                        min.node.size = bst$bestTune[[3]],
                        importance = "impurity")

barplot(sort(ranger::importance(gbm.final.imp), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("cyan", "blue"))(19))
```



Boosting - Test Error

```
pred.gbm <- predict(bst, newdata = test)
```

```
## Using 1235 trees...
```

```
RMSE(pred.gbm, y_test)
```

```
## [1] 1717.951
```

Among predictors, the most significantly one influences the model's ability to predict the 'Outstate' variable. In both measures, 'Expend' emerges as the most influential variable, suggesting that expenditure per student is a dominant predictor. This is followed by academic-related factors such as 'Terminal', 'PhD', and student performance metrics 'Top10perc' which also hold significant importance, reflecting the relevance of academic excellence and resources in predicting 'Outstate'.

The test error is 1717.951023.

Problem 2

```

auto = read.csv("/Users/zeze/Library/Mobile Documents/com~apple~CloudDocs/2024/24S BIST P8106 DS II/hw4
auto = auto |>
  drop_na() |>
  mutate(mpg_cat = as.factor(mpg_cat),
         origin = as.factor(origin))
head(auto)

```

```

##   cylinders displacement horsepower weight acceleration year origin mpg_cat
## 1         8           307         130   3504          12.0   70        1     low
## 2         8           350         165   3693          11.5   70        1     low
## 3         8           318         150   3436          11.0   70        1     low
## 4         8           304         150   3433          12.0   70        1     low
## 5         8           302         140   3449          10.5   70        1     low
## 6         8           429         198   4341          10.0   70        1     low

```

```

data_split <- initial_split(auto, prop = 0.7)
train2 <- training(data_split)
test2 <- testing(data_split)
#indexTrain2 <- createDataPartition(y = auto$mpg_cat, p = 0.7, list = FALSE)
#train2 <- auto[indexTrain2, ]
#test2 <- auto[-indexTrain2, ]
head(train2)

```

```

##   cylinders displacement horsepower weight acceleration year origin mpg_cat
## 1         4           90          48   2085          21.7   80        2    high
## 2         4          140          83   2639          17.0   75        1    high
## 3         4          122          80   2451          16.5   74        1    high
## 4         8          260          90   3420          22.2   79        1    high
## 5         4          156         105   2745          16.7   78        1    high
## 6         8          318         150   4190          13.0   76        1     low

```

(a) Build a classification tree using the training data, with mpg cat as the response. Which tree size corresponds to the lowest cross-validation error? Is this the same as the tree size obtained using the 1 SE rule?

Classification Tree

```

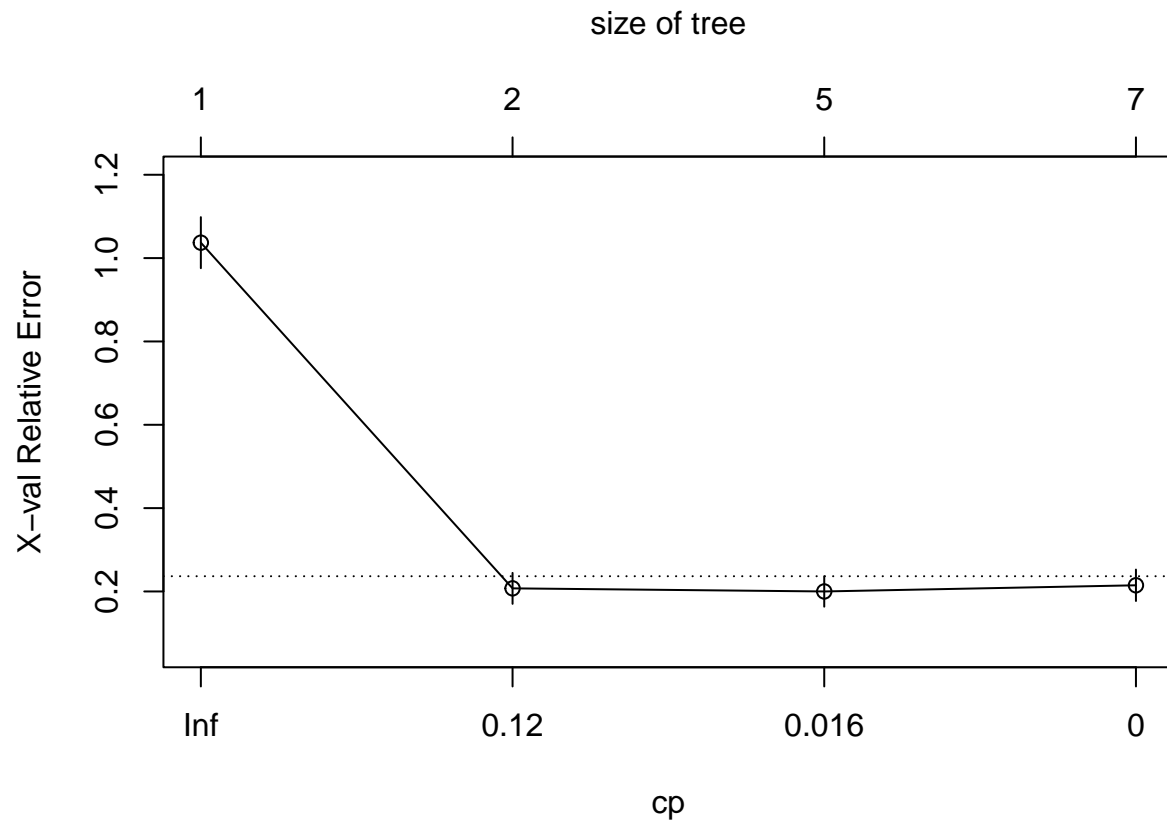
set.seed(1)
tree1 <- rpart(formula = mpg_cat ~ . , data = train2,
               control = rpart.control(cp=0))
cpTable <- printcp(tree1)

##
## Classification tree:
## rpart(formula = mpg_cat ~ . , data = train2, control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] displacement horsepower weight year
##

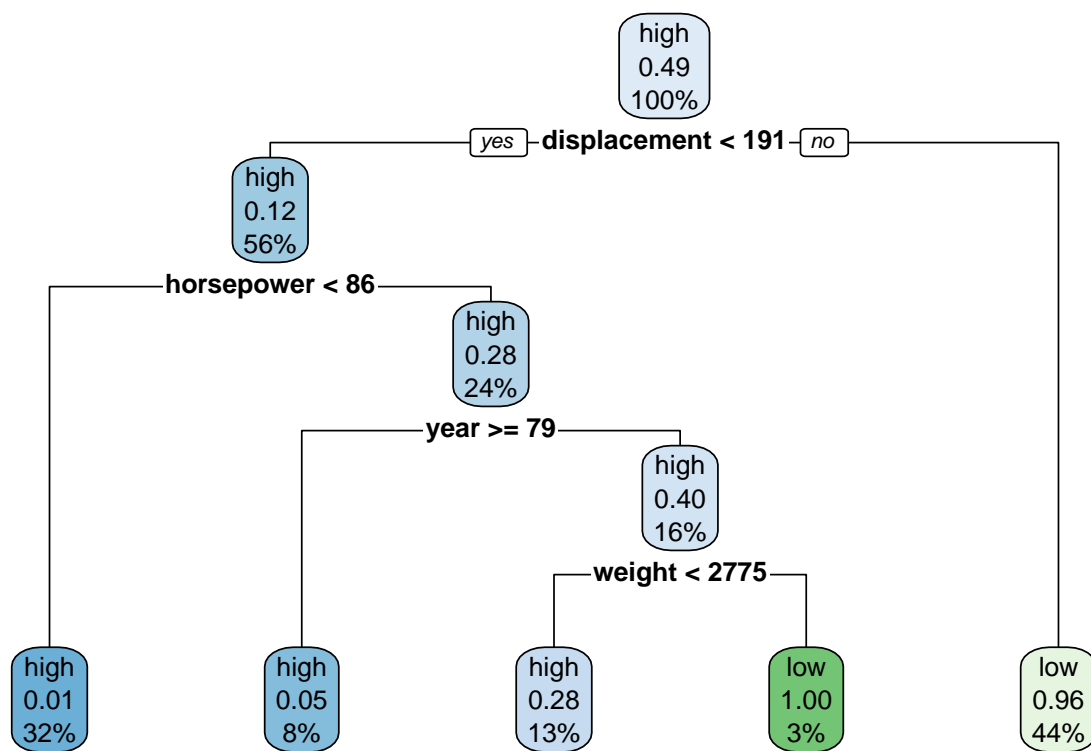
```

```
## Root node error: 135/274 = 0.4927
##
## n= 274
##
##      CP nsplit rel error  xerror   xstd
## 1 0.822222      0 1.000000 1.03704 0.061293
## 2 0.017284      1 0.177778 0.20741 0.037140
## 3 0.014815      4 0.125926 0.20000 0.036544
## 4 0.000000      6 0.096296 0.21481 0.037720
```

```
plotcp(tree1)
```



```
minErr <- which.min(cpTable[,4])
tree2 <- rpart::prune(tree1, cp = cpTable[minErr,1])
rpart.plot(tree2)
```

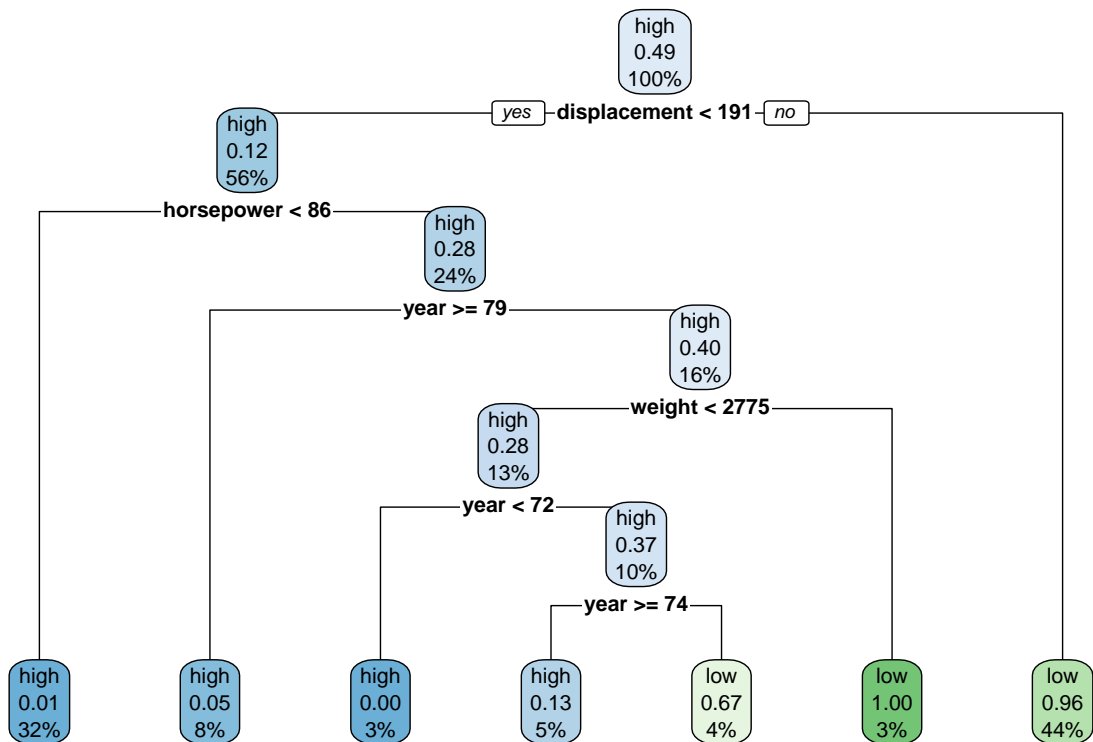


1se

```

minErr <- which.min(cpTable[, "xerror"])
minCVError <- cpTable[minErr, "xerror"]
minErrSE <- cpTable[minErr, "xstd"]
seIndex <- max(which(cpTable[, "xerror"] <= (minCVError + minErrSE)))
tree3 <- rpart::prune(tree1, cp = cpTable[seIndex, "CP"])
rpart.plot(tree3)

```



The tree size corresponds to the lowest cross-validation error is different after applying 1se.

(b) Perform boosting on the training data and report the variable importance. Report the test data performance.

Boosting

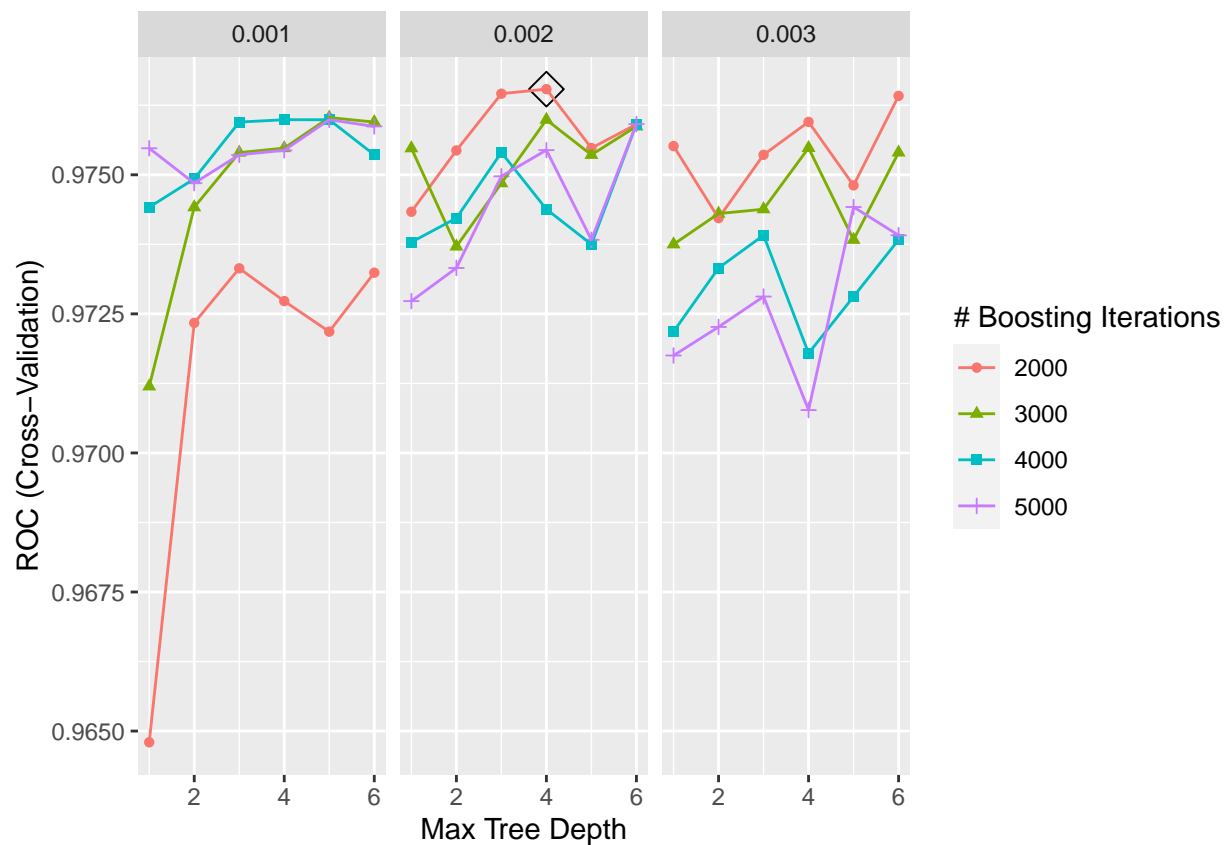
```

ctrl=trainControl(method = "cv",
                  classProbs = TRUE,
                  summaryFunction = twoClassSummary)

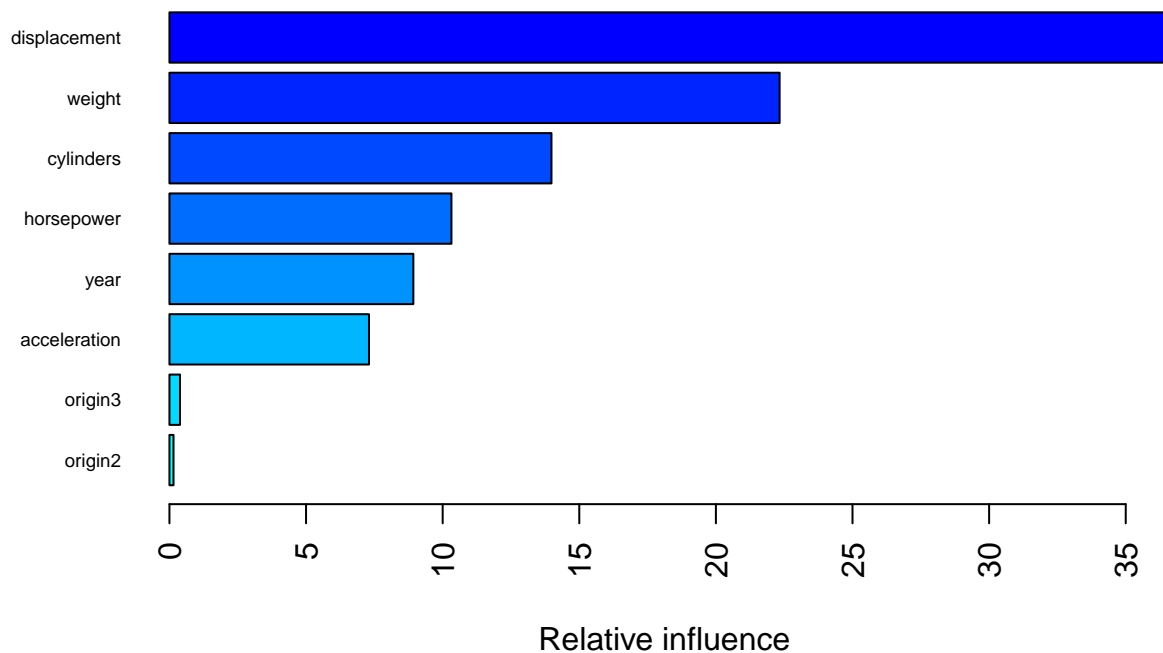
set.seed(1)
gbmA.grid <- expand.grid(n.trees = c(2000,3000,4000,5000),
                        interaction.depth = 1:6,
                        shrinkage = c(0.001, 0.002, 0.003),
                        n.minobsinnode = 1)

set.seed(1)
gbmA.fit <- train(mpg_cat ~ . , train2,
                 tuneGrid = gbmA.grid,
                 trControl = ctrl,
                 method = "gbm",
                 distribution = "adaboost",
                 metric = "ROC",
                 verbose = FALSE)
ggplot(gbmA.fit, highlight = TRUE)

```



```
gbmA.pred <- predict(gbmA.fit, newdata = test2, type = "prob")[,1]
summary(gbmA.fit$finalModel, las = 2, cBars = 19, cex.names = 0.6)
```



```
##           var    rel.inf
## displacement displacement 36.5969279
## weight          weight 22.3328029
## cylinders        cylinders 13.9822973
## horsepower      horsepower 10.3212955
## year            year 8.9275619
## acceleration acceleration 7.3034086
## origin3         origin3 0.3864974
## origin2         origin2 0.1492085
```

From the plot, we can see that “Displacement” appears to be the most influential variable, followed by “weight.” The variables “origin2” and “origin3” have no bar extending to the right, indicating they have zero or negligible importance in this context. The purpose of the model is not specified, but given the variables, it may be related to vehicles or engines.

Test Performance

```
gbmA.probs <- predict(gbmA.fit, newdata = test2, type = "prob")
roc(response = test2$mpg_cat, predictor = gbmA.probs[, "high"])
```

```
## Setting levels: control = high, case = low
```

```
## Setting direction: controls > cases
```



```
##  
## Call:  
## roc.default(response = test2$mpg_cat, predictor = gbmA.probs[,      "high"])  
##  
## Data: gbmA.probs[, "high"] in 57 controls (test2$mpg_cat high) > 61 cases (test2$mpg_cat low).  
## Area under the curve: 0.9888
```

The AUC value is 0.99, which is very close to 1. This indicates an excellent performance of the model on the test data, with high accuracy in differentiating between the ‘high’ and ‘low’ categories of the ‘mpg_cat’ variable. The ‘controls’ are instances labeled as ‘high’ and ‘cases’ as ‘low’. An AUC value above 0.9 is typically considered outstanding, suggesting that the model’s predicted probabilities (gbmA.probs[, “high”]) are highly effective at ranking the test data instances with a high degree of separation between the two mpg categories.