

Dimension Reduction Methods in Linear Regression

Yifei Sun, Runze Cui

Contents

SVD	2
Principal components regression (PCR)	3
Partial least squares (PLS)	4
PCR and PLS using caret	6
PCR	6
PLS	7

```
library(ISLR)
library(pls)
library(caret)
library(tidymodels)
```

Predict a baseball player's salary on the basis of various statistics associated with performance in the previous year. Use `?Hitters` for more details.

```
data(Hitters)
Hitters <- na.omit(Hitters)
set.seed(2222)

data_split <- initial_split(Hitters, prop = 0.8)

# Extract the training and test data
training_data <- training(data_split)
testing_data <- testing(data_split)

# training data
x <- model.matrix(Salary ~ ., training_data)[, -1]
y <- training_data$Salary

# test data
x2 <- model.matrix(Salary ~ ., testing_data)[, -1]
y2 <- testing_data$Salary
```

SVD

```
# center and scale
x3 <- scale(x)

# SVD
x_svd <- svd(x3)
u <- x_svd$u
v <- x_svd$v
d <- diag(x_svd$d)
# corplot::corplot(t(u) %*% u, is.corr = FALSE)
# corplot::corplot(t(v) %*% v, is.corr = FALSE)
# corplot::corplot(v %*% t(v), is.corr = FALSE)
# corplot::corplot(d, is.corr = FALSE)

# definition
x4 <- u %*% d %*% t(v)
all.equal(x3, x4, check.attributes = FALSE)

## [1] TRUE

# PCA
x_pca <- prcomp(x, scale. = TRUE)

all.equal(x_pca$rotation, v, check.attributes = FALSE)

## [1] TRUE
```

Principal components regression (PCR)

We fit the PCR model using the function `pcr()`.

```
set.seed(2)
pcr.mod <- pcr(Salary ~ .,
               data = training_data,
               scale = TRUE, # scale = FALSE by default
               validation = "CV")

summary(pcr.mod)
```

```
## Data:      X dimension: 210 19
## Y dimension: 210 1
## Fit method: svdpc
## Number of components considered: 19
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV           472.1   371.4   372.6   374.2   373.3   369.3   368.9
## adjCV        472.1   370.9   371.9   373.4   372.4   368.3   367.7
##      7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV       368.6   369.6   370.0   367.4   368.3   369.1   377.5
## adjCV    367.4   368.3   368.6   365.7   366.6   367.4   375.4
##      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## CV       368.8   369.9   358.4   358.2   356.1   360.6
## adjCV    366.5   367.6   356.1   355.9   353.7   357.9
##
## TRAINING: % variance explained
##      1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8 comps
## X         37.90   59.83   71.09   79.07   84.27   88.44   92.14   94.77
## Salary    40.56   41.42   41.56   42.62   44.55   45.57   46.06   46.06
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X         96.12   97.10   97.89   98.60   99.12   99.47   99.74
## Salary    46.40   48.56   48.60   48.61   48.66   51.26   51.33
##      16 comps 17 comps 18 comps 19 comps
## X         99.89   99.97   99.99   100.00
## Salary    54.55   55.02   55.93   56.02
```

```
validationplot(pcr.mod, val.type = "MSEP", legendpos = "topright")
```



```
cv.mse <- RMSEP(pcr.mod)
ncomp.cv <- which.min(cv.mse$val[1,,]) - 1
ncomp.cv

## 18 comps
##      18

predy2.pcr <- predict(pcr.mod, newdata = testing_data,
                      ncomp = ncomp.cv)

# test MSE
mean((y2 - predy2.pcr)^2)

## [1] 90894.84
```

Partial least squares (PLS)

We fit the PLS model using the function `pls()`.


```
set.seed(2)
pls.mod <- pls(Salary~.,
               data = training_data,
               scale = TRUE,
               validation = "CV")

summary(pls.mod)

## Data:      X dimension: 210 19
## Y dimension: 210 1
## Fit method: kernelppls
## Number of components considered: 19
##
```

```
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           472.1   369.1   369.7   369.6   365.7   367.5   366.2
## adjCV        472.1   368.5   368.5   368.0   364.2   365.4   363.8
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV           363.7   360.7   360.8   359.6   355.5   356.1   361.3
## adjCV        361.4   358.3   358.4   357.1   353.1   353.9   358.5
##      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## CV           358.0   356.4   357.7   357.8   357.8   360.6
## adjCV        355.5   354.1   355.2   355.3   355.3   357.9
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X           37.70   50.45   65.18   71.93   77.45   84.28   88.12   89.70
## Salary      42.77   46.29   47.94   49.63   51.55   52.45   53.26   54.57
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X           92.76   94.43   95.31   97.55   98.02   98.40   98.92
## Salary      54.80   55.11   55.44   55.50   55.68   55.88   55.92
##      16 comps 17 comps 18 comps 19 comps
## X           99.13   99.64   99.99  100.00
## Salary      55.97   55.98   55.99   56.02
```

```
validationplot(pls.mod, val.type = "MSEP", legendpos = "topright")
```



```
cv.mse <- RMSEP(pls.mod)
ncomp.cv <- which.min(cv.mse$val[1,,]) - 1
ncomp.cv
```

```
## 11 comps
##      11
```

```

predy2.pls <- predict(pls.mod, newdata = testing_data,
                     ncomp = ncomp.cv)
# test MSE
mean((y2 - predy2.pls)^2)

## [1] 92933.01

```

PCR and PLS using caret

PCR

```

ctrl1 <- trainControl(method = "repeatedcv",
                     number = 10,
                     repeats = 5,
                     selectionFunction = "best") # "oneSE" for the 1SE rule

# show information about the model
modelLookup("pcr")

##   model parameter      label forReg forClass probModel
## 1   pcr      ncomp #Components   TRUE   FALSE   FALSE
modelLookup("pls")

##   model parameter      label forReg forClass probModel
## 1   pls      ncomp #Components   TRUE   TRUE   TRUE

# Two ways for standardizing predictors

# train(..., preProc = c("center", "scale"))
set.seed(2)
pcr.fit <- train(x, y,
                method = "pcr",
                tuneGrid = data.frame(ncomp = 1:19),
                trControl = ctrl1,
                preProcess = c("center", "scale"))

predy2.pcr2 <- predict(pcr.fit, newdata = x2)
mean((y2 - predy2.pcr2)^2)

## [1] 93174.43


# pcr(..., scale = TRUE)
set.seed(2)
pcr.fit2 <- train(x, y,
                method = "pcr",
                tuneGrid = data.frame(ncomp = 1:19),
                trControl = ctrl1,
                scale = TRUE)

predy2.pcr3 <- predict(pcr.fit2, newdata = x2)
mean((y2 - predy2.pcr3)^2)

## [1] 93174.43

```

```
ggplot(pcr.fit, highlight = TRUE) + theme_bw()
```




PLS

```
set.seed(2)
pls.fit <- train(x, y,
  method = "pls",
  tuneGrid = data.frame(ncomp = 1:19),
  trControl = ctrl1,
  preProcess = c("center", "scale"))
predy2.pls2 <- predict(pls.fit, newdata = x2)
mean((y2 - predy2.pls2)^2)
```

```
## [1] 92933.01
```

```
ggplot(pls.fit, highlight = TRUE)
```



Here are some old code on elastic net.

```
set.seed(2)
enet.fit <- train(Salary ~ .,
                  data = training_data,
                  method = "glmnet",
                  tuneGrid = expand.grid(alpha = seq(0, 1, length = 21),
                                         lambda = exp(seq(6, 0, length = 100))),
                  trControl = ctrl1)

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.

# myCol <- rainbow(25)
# myPar <- list(superpose.symbol = list(col = myCol),
#               superpose.line = list(col = myCol))
# plot(enet.fit, xTrans = log, par.settings = myPar)
```

Comparing the models based on resampling results.

```
resamp <- resamples(list(elastic_net = enet.fit,
                        pcr = pcr.fit,
                        pls = pls.fit))

summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
```



```
## Models: elastic_net, pcr, pls
## Number of resamples: 50
##
## MAE
##           Min. 1st Qu.  Median    Mean 3rd Qu.    Max. NA's
## elastic_net 163.3124 215.5571 240.4248 248.2230 275.3263 373.5337    0
## pcr         171.4586 222.0582 243.1119 253.4684 275.8308 379.9711    0
## pls         164.1584 217.9899 241.4971 251.3777 275.9601 379.1324    0
##
## RMSE
##           Min. 1st Qu.  Median    Mean 3rd Qu.    Max. NA's
## elastic_net 220.0769 289.7156 325.0586 344.2359 371.1565 572.1447    0
## pcr         222.8634 286.5283 334.7599 345.7586 372.2937 563.5358    0
## pls         220.7667 286.4764 323.6777 344.4050 380.5525 554.8279    0
##
## Rsquared
##           Min. 1st Qu.  Median    Mean 3rd Qu.    Max. NA's
## elastic_net 0.01505160 0.3869746 0.4944062 0.4811950 0.6202539 0.8622079    0
## pcr         0.02684558 0.3883547 0.4925494 0.4750800 0.6026677 0.8347328    0
## pls         0.02822971 0.3507958 0.5037413 0.4775314 0.6066276 0.8513646    0
```

```
bwplot(resamp, metric = "RMSE")
```

